

✓ **Congratulations! You passed!**

Next Item



1. What are some of the key goals of the estimator API?

1 / 1
point

- ☐ Create production-ready machine learning models using an API
- ☐ Train on large datasets that do not fit in memory
- ☐ Quickly monitor your training metrics in Tensorboard
- ☒ All of the above

Correct



2. What is one of the largest benefits of the estimator API?

1 / 1
point

- ☐ It automatically tunes your ML model hyperparameters for you
- ☐ It requires you to specify which hardware you will run on for the best performance
- ☒ It abstracts away boilerplate code which saves you time

Correct



3. What is the right way to call a linear regression model with tf.estimator?

1 / 1
point

- ☐ tf.estimator.line_model
- ☐ tf.estimator.regression
- ☒ tf.estimator.LinearRegressor
- ☐ tf.estimator.LinearClassifier

Correct



4. Inputs to the estimator model are in the form of:

1 / 1
point

- ☐ BigQuery datasets
- ☐ scalars
- ☐ hyperparameters
- ☒ feature columns

Correct



5. Numeric inputs can be passed to a linear regressor as-is, but categorical columns are often:

1 / 1
point

- ☐ Not used, only numeric values can be passed
- ☐ Uniformly distributed and aggregated first
- ☒ One-hot encoded
- ☐ Cleansed because of duplicate records

Correct



6. What is the size of the training dataset (features + labels) in this example?

1 / 1
point

```
def train_input_fn():
    features = {"sq_footage": [ 1000, 2000, 3000, 1000, 2000, 3000],
               "type": ["house", "house", "house", "apt", "apt", "apt"]}
    # prices in thousands
    labels = [ 500, 1000, 1500, 700, 1300, 1900]
    return features, labels
```

- ☐ 6 rows, 2 columns
- ☐ 7 rows, 4 columns
- ☒ 6 rows, 3 columns

Correct


- ☐ 6 rows, 4 columns

✓

7. In this example, what extra parameters does the DNNRegressor take that the LinearRegressor doesn't?

1 / 1
point

```
model = tf.estimator.DNNRegressor(featcols,
                                  hidden_units=[3, 2])
```



tf.estimator.
LinearRegressor

- ☒ hidden_units

Correct

- ☐ featcols
- ☐ regression
- ☐ neurons

✓

8. In what situation do you have to delete the model directory before starting training?

1 / 1
point

- ☒ If you have changed the model structure from the previous time, for example, you used a DNNRegressor with [64,32] last time and now you are using [32, 16]

Correct

The old checkpoints are no longer valid for your new model structure. So, you have to start afresh

- ☐ If your model is not performing well enough and you need to train for more epochs or with additional examples
- ☐ If you want to automatically checkpoint from an earlier saved model

✓

9. What is the difference between steps and max_steps?

1 / 1
point

```
def pandas_train_input_fn(df): # a Pandas dataframe
    return tf.estimator.inputs.pandas_input_fn(
        x = df,
        y = df['price'],
        batch_size=128,
        num_epochs=10,
        shuffle=True
    )

model.train(pandas_train_input_fn(df))
model.train(pandas_train_input_fn(df), steps=1000)
model.train(pandas_train_input_fn(df), max_steps=1000)
```

Trains until input exhausted (10 epochs) starting from checkpoint

1000 additional steps from checkpoint

1000 steps - might be nothing if checkpoint already there

☒ Steps means "train these many additional steps". max_steps means "train up to these many steps total, starting from how many ever steps have been completed so far"

Correct

- ☐ There is no difference
- ☐ Steps means "train this many steps total". max_steps means "train these many additional steps"

👤 🗨️ 📄